

AURA Forge Stakeholder One-Pager

Governed adaptive engineering for FitGPT

Problem

FitGPT is a real full-stack application, so engineering changes vary widely in risk. A fixed AI workflow treats a small documentation update too much like a governance-sensitive code change, wasting agent time and human approval effort.

What AURA Forge Does

AURA Forge classifies each change as LOW, MEDIUM, or HIGH using deterministic rules, then routes it through the minimum justified mix of agents, deterministic checks, policy gates, and human approvals.

Why It Matters

The goal is faster routine engineering without weakening governance. Low-risk changes avoid unnecessary overhead. Sensitive changes keep full controls.

Measured Results

Scenario	Quality	Cost	Roles	Humans
LOW	14/16 -> 16/16	-44.32%	-60%	-100%
MEDIUM	15/16 -> 16/16	-19.71%	-40%	-50%
HIGH	15/16 -> 16/16	-5.89%	unchanged	unchanged
Aggregate	44/48 -> 48/48	-19.22%	-33.33%	-50%

Risks Managed

Least-privilege MCP boundaries, real overreach denial, policy tests, pipeline-integrity checks, audit artifacts, CI gates, and Change Passport evidence.

Limitations

Measured results cover three representative scenarios only. No production deployment, real-user data, or company-wide savings claim is made. HIGH work intentionally keeps governance.

Next Step

Use AURA Forge as the course-proven paved road for future isolated FitGPT engineering experiments. Do not deploy it to production without a separate production-readiness review.