

AURA Forge Ops-Ready Runbook

Production-like capstone operating procedure, not a live production on-call runbook

This runbook covers the isolated AURA Forge capstone workflow. It defines monitored signals, failure response, escalation, rollback, deterministic fallback, recovery verification, and evidence preservation without implying a real production on-call environment.

1. Monitoring Signals

Signal	Source	Action threshold
Policy gate failure	GitHub Actions policy-tests	Stop. Do not merge until diagnosed and fixed.
Evaluation gate failure	GitHub Actions evaluation-gate	Stop. Separate workflow/dependency failure from real regression.
Pipeline integrity failure	GitHub Actions pipeline-integrity	Stop. Treat as guardrail weakening until proven otherwise.
Advisory unavailable	GitHub Actions advisory-review	Continue only if permanent gates pass and skip artifact is explicit.
Audit trail failure	GitHub Actions audit-trail	Stop for evidence repair unless an upstream permanent failure explains missing artifacts.
Role overreach denial	MCP authorization result	Expected for unauthorized attempts; preserve as safety evidence.
Holdout checksum mismatch	shasum docs/holdout-task-set.md	Stop. Restore integrity without rewriting history.
Secret-pattern finding	Local scan or GitHub secret scanning	Stop. Do not push. Sanitize and rotate if real secret was exposed.

2. Incident Procedure

- Classify the incident: policy/eval/integrity, advisory-only, audit/evidence, environment/dependency, or security.
- Preserve failed run ID, commit SHA, failing job, and sanitized log excerpt.
- Reproduce locally where possible without production access.
- Apply the smallest evidence-backed fix.
- Run local deterministic verification.
- Commit normally on capstone/aura-forge; do not force-push.
- Push to origin/capstone/aura-forge and confirm the next GitHub run passes permanent gates.

3. Failure Response Matrix

Failure	Immediate action	Escalation / rollback	Recovery verification
Classifier failure	Stop routing; treat malformed evidence as HIGH.	Revert classifier rule or route to HIGH with escalation.	Risk classifier and CI change-classifier tests.
Routing failure	Stop; do not execute malformed plan.	Revert router change or require plan approval.	Adaptive router tests.
Policy gate failure	Stop; do not mark ready.	Revert permission widening or repair allow-list/test mismatch.	Policy/MCP/coursetools tests.
Evaluation failure	Stop; determine real regression vs environment issue.	Revert real regression or repair workflow dependency.	Focused pytest plus green CI evaluation gate.
MCP denial	Preserve denial and confirm no external state changed.	If unexpected, escalate to HIGH and inspect allow-list.	Governance-overreach and policy tests.
CI environment failure	Preserve run and exact missing prerequisite.	Normal fix commit; no force-push.	Next GitHub Actions run passes.
Advisory AI unavailable	Continue only if permanent gates pass.	No model-review fabrication.	Skip artifact plus audit summary.

Failure	Immediate action	Escalation / rollback	Recovery verification
Change Passport producer missing	Stop; do not manually fill fields or fabricate readiness.	Regenerate from real producers or leave absent with explanation.	Change Passport test and generator.
Secret finding	Stop and do not push.	Sanitize; rotate if exposure happened.	Secret scan returns zero findings.

4. Deterministic Fallback And Graceful Degradation

- Unknown or malformed classifier/router evidence fails closed to HIGH or escalation.
- Advisory AI is optional; missing credentials produce a sanitized SKIPPED - AI SECRET UNAVAILABLE artifact.
- Change Passport producer missing or other evidence failure stops readiness rather than inventing fields.
- No production service, production database, production secret, or live user is required for deterministic gates.

5. Evidence To Preserve

- GitHub run ID, commit SHA, job name, and sanitized logs.
- Classifier output, route plan, changed paths, and gate outputs.
- Policy reports, evaluation reports, integrity reports, advisory skip artifacts, and audit trail.
- Governance denial evidence including role, tool, decision, reason, and external-state flag.
- Holdout checksum evidence and secret-scan results.

Source Evidence

Primary sources: docs/capstone/monitoring-incident-runbook.md, docs/capstone/governance-ci-results.md, docs/capstone/tool-evolution-drill.md, eval/test_adaptive_router.py, scripts/check-pipeline-integrity.py.