

AURA Forge ADR Package

Combined architecture decision records for Canvas upload

Table of Contents

ADR	Decision area
0001	Quality rubric and evaluation design
0002	Memory, context, and prompt architecture
0003	MCP tool boundaries and least privilege
0004	Adaptive risk routing
0005	Governance and eval-gated CI
0006	Agent-to-deterministic conversion
0007	Change Passport evidence design

ADR 0001: Quality Rubric And Evaluation Design

Context

Comparable evidence was needed across PRE-AURA and AURA runs.

Decision

Use a 16-point rubric across outcome accuracy, route appropriateness, governance fidelity, and evidence/readiness quality.

Rejected Alternatives

- Single pass/fail result only.
- Cost-only comparison.

Evidence

- PRE-AURA LOW: 14/16 FAIL and route appropriateness 2/4.
- AURA aggregate: 44/48 -> 48/48.

Consequences

Supports route-quality claims without pretending cost is the only objective.

Open Risks

Calibrated to three representative scenarios.

ADR 0002: Memory, Context, And Prompt Architecture

Context

Role prompts, bounded evidence packets, and local memory rules were needed without making memory authoritative proof.

Decision

Keep role prompts narrow, pass required evidence only, and preserve role outputs as artifacts.

Rejected Alternatives

- Broad shared context for every role.
- Memory-only evidence.

Evidence

- Calibration preserved a Tester output-contract failure.
- AURA HIGH used separate role evidence with zero authorization denials.

Consequences

Repository artifacts remain authoritative and role evidence stays auditable.

Open Risks

Long runs still need final sanity checks to avoid stale-attempt mixing.

ADR 0003: MCP Tool Boundaries And Least Privilege

Context

Governance depends on role-scoped MCP grants and real authorization evidence.

Decision

Keep role/tool boundaries explicit and testable.

Rejected Alternatives

- Prompt-only boundaries.
- Broad shell or task-tracker access for implementation roles.

Evidence

- Permanent policy/MCP/coursetools checks.
- GO-20260811-001 denied implementer access to task_tracker.

Consequences

Least privilege can be shown with machine evidence.

Open Risks

The task_tracker is a course simulator, not a live external tracker.

ADR 0004: Adaptive Risk Routing

Context

The fixed PRE-AURA route over-served LOW and did not adapt to risk.

Decision

Use deterministic aura-risk-v1 classification and frozen route plans per tier.

Rejected Alternatives

- Model-decided risk routing.
- Always full route.
- Always minimize cost.

Evidence

- LOW roles 5 -> 2 and checkpoints 2 -> 0.
- HIGH retained 5 roles and 2 checkpoints.

Consequences

Route choice is reproducible and risk-sensitive.

Open Risks

Future HIGH categories may need additional gates.

ADR 0005: Governance And Eval-Gated CI

Context

The course needed a capstone governance workflow beyond ordinary app tests.

Decision

Add CI jobs for classification, policy tests, conditional evaluation, pipeline integrity, advisory review, and audit trail.

Rejected Alternatives

- Replace existing app tests.
- Make AI review blocking.
- Skip audit on failure.

Evidence

- Pipeline checker rejects weak policy gates and broad permissions.
- Governed Docker checks reported 18 passed.

Consequences

Permanent gates work without optional AI secrets.

Open Risks

Runner deprecation warnings are external and non-blocking.

ADR 0006: Agent-To-Deterministic Conversion

Context

One stable docs/config fact did not need repeated agentic review.

Decision

Convert DATABASE_URL representation checking to deterministic code.

Rejected Alternatives

- Convert the whole auditor.
- Call the original risk classifier the conversion.

Evidence

- Script emits JSON with \$0 model cost and about 0.156 ms local runtime.
- Pass/fail tests cover behavior.

Consequences

Stable factual guardrails are cheaper, faster, and easier to audit.

Open Risks

Covers one narrow fact, not the entire backend config surface.

ADR 0007: Change Passport Evidence Design

Context

A terminal summary was not enough to make readiness evidence portable and machine-checkable.

Decision

Build Change Passport as deterministic evidence aggregation from real producers.

Rejected Alternatives

- Manually filled passport fields.
- Add CI fields before real runs.
- Make Passport a router.

Evidence

- Builder reads classifier/router, metrics, quality, approvals, policy tests, CI fields, and hashes.
- Missing producers fail tests.

Consequences

Readiness becomes portable and machine-checkable.

Open Risks

CI fields should only refresh from real GitHub producers.

Source Evidence

Primary sources: docs/adr/0001 through 0007 and docs/capstone/adr-evidence-matrix.md. This package preserves the ADR context, decision, rejected alternatives, evidence, consequences, and open risks without rewriting ADR history.