

AURA Forge Rubric Self-Check

Current defensible score: 49 / 52

This self-check summarizes the 13 official criteria using current evidence. Scores are not inflated. Criteria 2, 4, 9, and 11 remain conservative at 3/4.

Criterion	Score	Strong evidence	Short reason
1. Workflow Scoping	4/4	docs/capstone/aura-forge-prd.md; representative-scenarios.md; control-vs-aura-impact.md	Quantified baseline, acceptance criteria, and why adaptive governance is the right tool.
2. Sandboxed Environment	3/4	Dockerfile; README.md; reproducibility-runbook.md	Read-only verification paths and runnable guidance exist, but network egress is not proven airtight.
3. Quality Spec & Baseline	4/4	quality-rubric.md; control-baseline-results.md; ADR 0001	16-point rubric and precise PRE-AURA baseline support future comparison.
4. Agents, Skills & Memory	3/4	.claude/agents; .claude/skills; docs/memory-architecture.md	Role/memory structure exists, but complete stale-memory pruning and before/after reflection scores are not fully reconstructable.
5. Orchestration & MCP Tools	4/4	routing-and-tool-grant-map.md; mcp/coursetools_server.py; eval/test_policy.py	Role/tool map, narrow grants, and least-privilege enforcement are concrete.
6. Evaluation & Calibration	4/4	eval/*; docs/calibration-log.md; docs/holdout-task-set.md	Normal, edge, adversarial, regression, policy, and holdout evidence exist.
7. Governance, Security & CI/CD	4/4	governance-overreach-demo.md; governance-ci-results.md; scripts/check-pipeline-integrity.py	Overreach denial, red-team/policy checks, CI gates, and sanitized evidence are present.
8. Right-Tool Decisions & ADRs	4/4	right-tool-decision-matrix.md; adr-evidence-matrix.md; docs/adr/*	Tool-choice matrix, quantitative deterministic conversion, and evidence-backed ADRs exist.
9. Production Integration & Tool-Evolution Drill	3/4	tool-evolution-drill.md; monitoring-incident-runbook.md; eval/test_adaptive_router.py	Production-like CI and drill exist, but the fault was intentional, not an unobserved spontaneous regression.
10. Iteration Narrative & Impact	4/4	iteration-narrative.md; control-vs-aura-impact.md; calibration-log.md	Coherent engineering story and bounded measured impact are backed by logs/evidence.
11. Stakeholder Communication	3/4	stakeholder-one-pager.md; final Gamma deck; https://youtu.be/WxEXCWC75vs	Video artifact exists, but no independent local transcript/content review proves every 4/4 video requirement.
12. Clarity & Flow	4/4	AURA_Forge_Final_Gamma_Presentation.pdf; teleprompter; video staging	Final deck and recording sequence provide clear technical-to-stakeholder flow.
13. Design	4/4	AURA_Forge_Final_Gamma_Presentation.pdf; submission assets	Visual deck supports architecture, governance, CI, Passport, and measured-result storytelling.

Score Summary

Current defensible score: 49 / 52. The score remains conservative. It does not claim airtight sandboxing, complete historical memory-pruning evidence, an unobserved spontaneous tool-evolution regression, or independently verified 4/4 video content.

Source Evidence

Primary source: docs/capstone/final-rubric-audit.md, with supporting evidence in docs/capstone/evidence-index.md and docs/capstone/official-rubric-reference.md.