

AURA Forge Rubric Self-Check

Official LaunchCode 13-criterion rubric · corrected non-video revision · 2026-09-01

48 / 52

DEFENSIBLE
SELF-ASSESSMENT

Score calculation: 52 – 1 (criterion 2) – 1 (criterion 4) – 1 (criterion 9) – 1 (criterion 11) = 48. This is a self-assessment, not a Canvas regrade.

Criteria 1–7

#	Official criterion	Score	Defensible evidence summary
1	Workflow Scoping	4/4	Quantified PRE-AURA comparison, measurable acceptance criteria, and explicit comparison with generic-agent, fixed-automation, and uniform-route alternatives.
2	Sandboxed Environment	3/4	Fresh offline verifier passed 32 tests with no network, read-only root/workspace, no credentials, dropped capabilities, bounded resources, and unique run naming. Historical model-backed runs used bridge networking.
3	Quality Spec & Baseline	4/4	Complete PRD, four-dimension scoring guide with examples and thresholds, plus quality, review-latency proxy, defect-rate proxy, cycle-time proxy, and cost.
4	Agent, Skills & Memory	3/4	Versioned agents and three skills, scoped memory, and before/fix/rerun reflection evidence. No stale memory entry was eligible for a real pruning event.
5	Orchestration & MCP Tools	4/4	Scoped multi-agent route, least-privilege grants, persistent storage, and deterministic semantic-vector retrieval with schema, classification ceilings, and citations.
6	Evaluation & Calibration	4/4	Deterministic and rubric harnesses, tracked holdout, representative runs, red-team evidence, and evidence-linked calibration records.
7	Governance, Security & CI/CD	4/4	Real overreach denial, enforced CI policy/eval/integrity gates, red-team prompts, audit trail, rollback criteria, and sanitized artifacts.

AURA Forge Rubric Self-Check

Official LaunchCode 13-criterion rubric · corrected non-video revision · 2026-09-01

48 / 52

DEFENSIBLE
SELF-ASSESSMENT

Criteria 8–13

#	Official criterion	Score	Defensible evidence summary
8	Right-Tool Decisions & ADRs	4/4	Reusable agent/deterministic/human matrix, nine evidence-backed ADRs, and quantified deterministic conversion with explicit comparison limitations.
9	Production Integration & Tool-Evolution Drill	3/4	Timeout/retry/budget decisions and rollback/escalation are tested. The preserved tool-evolution fault was intentional, so it is not misrepresented as an unobserved spontaneous regression.
10	Iteration Narrative & Impact	4/4	Multiple before/fix/rerun cycles and a five-metric impact comparison with measured/proxy labels and calculations.
11	Stakeholder Communication	3/4	One-pager translates percentages into measured dollars and time. Canvas reports the current video is about 13 minutes against a 10-minute cap, so video work remains open.
12	Clarity & Flow	4/4	Final Gamma deck and speaking materials preserve a coherent problem-to-evidence narrative.
13	Design	4/4	Final deck uses consistent hierarchy and evidence visuals for routing, governance, CI, and measured impact.

Current verification

Full local evaluation suite: 114 passed · Offline network-disabled container: 32 passed · GitHub CI 33517133584: SUCCESS · Public anonymous access: HTTP 200 · Video unchanged.

Required human actions before resubmission

- Submit the exact public repository link through Canvas/LaunchCode.
- Confirm the grader received and can open the full repository.
- Complete the walkthrough at or below 10 minutes with Muhammad.

Repository delivery is ready: implementation commit a7813c4 is public and passed governance CI. The remaining grading action is to submit the exact link and confirm receipt; the video-length correction remains separate.